

Lesson 3: Introduction to Packages

- What are R packages and why are they important?
- Installing packages from CRAN (`install.packages()`).
- Loading packages into your R session (`library()`, `require()`).
- Exploring package documentation.

While R's base installation is powerful, its true strength lies in its extensibility through **packages**. Packages are collections of functions, data, and compiled code in a well-defined format, designed to extend R's capabilities for specific tasks. They are the primary way the R community shares reusable code and functionalities.

Phase 1: Introduction to R and Fundamentals

Module 1.4: Functions and Packages

Lesson 3: Introduction to Packages

This lesson will introduce you to the concept of R packages, why they are indispensable for data analysis, and how to effectively manage them. You'll learn how to install, load, and find packages, as well as get an overview of some of the most commonly used and powerful packages in the R ecosystem.

1. What are R Packages and Why are They Used?

- **Definition:** An R package is a collection of **functions**, data, and documentation that extends the capabilities of R's base system. Think of them as add-ons or plugins for R.
- **Purpose:**
 - **Extend Functionality:** Provide specialized tools for statistics, machine learning, visualization, data manipulation, web scraping, and more.
 - **Code Reusability:** Prevent users from reinventing the wheel by providing pre-written, tested, and optimized code for common tasks.
 - **Community Contribution:** Allow R users and developers worldwide to share their work and contribute to the R ecosystem.
 - **Organization:** Group related functions and data together, making them easier to discover and use.

R's core functionality is powerful, but packages are what truly make R a versatile tool for almost any data-related task. The Comprehensive R Archive Network (CRAN) hosts over 20,000 packages, with thousands more available elsewhere (e.g., Bioconductor, GitHub).

2. Installing Packages (`install.packages()`) Before you can use a package, you need to install it onto your system. This typically only needs to be done once per package on a given R installation.

Syntax:

```
install.packages("package_name")
```

Parameters:

- **"package_name":** A character string specifying the name of the package you want to install. It **must** be enclosed in quotes.

Important Notes:

- You need an active internet connection to install packages from CRAN.

- You might be prompted to choose a CRAN mirror (a server from which to download). Choose one geographically close to you for faster downloads.
- `install.packages()` downloads the package from CRAN (or another specified repository) and places it in your R library folder.

Code Snippets:

```
print("--- Installing Packages ---")

## [1] "--- Installing Packages ---"
# Example: Installing a popular data manipulation package, 'dplyr'
# If you've installed it before, R will usually just say it's already installed.
# It's good practice to check if it's installed first, though for a learning
# environment, simply running install.packages() is fine.

# You can check if a package is installed:
# if (!requireNamespace("dplyr", quietly = TRUE)) {
#   install.packages("dplyr")
# }

# Let's install a less common example for demonstration without affecting
# your existing setup too much, or use a package that is commonly used.
# 'stringr' is a good basic text manipulation package.

# Installing 'stringr'
# install.packages("stringr") # Uncomment to run

# To see where packages are installed (your R library paths):
print("Your R library paths:")

## [1] "Your R library paths:"
print(.libPaths())

## [1] "C:/Users/Titus/AppData/Local/R/win-library/4.4"
## [2] "C:/Program Files/R/R-4.4.3/library"
print("Installation successful (or package already installed) if no errors occurred.")

## [1] "Installation successful (or package already installed) if no errors occurred."
```

3. Loading Packages (`library()`, `require()`) Installing a package makes it available on your system, but to use its functions in your current R session, you must **load** it. This typically needs to be done at the beginning of each R session where you intend to use functions from that package.

- `library(package_name)`: The most common and recommended way to load a package. It loads the package into your current R session. If the package isn't found, it will produce an error and stop execution.
- `require(package_name)`: Similar to `library()`, but it returns `TRUE` if the package is successfully loaded and `FALSE` if not (without stopping execution). It's often used within functions or scripts where you want to conditionally load a package or check for its availability.

Syntax:

`library(package_name)` `require(package_name)` # Less common for interactive use

Parameters:

- `package_name`: The name of the package to load. **No quotes are generally needed** for `library()` or `require()`, although they work with quotes too.

Code Snippets:

```
print("--- Loading Packages ---")

## [1] "--- Loading Packages ---"

# Let's use the 'stringr' package which you might have installed in the previous step
# or is likely already installed on most R installations.

# Load the stringr package
library(stringr)
print("stringr package loaded using library().")

## [1] "stringr package loaded using library()."

# Now you can use functions from stringr, e.g., str_length()
text_data <- "Introduction to R Packages"
char_count <- str_length(text_data)
print(paste("Number of characters in '", text_data, "':", char_count))

## [1] "Number of characters in ' Introduction to R Packages ': 26"

# Using require()
# Usually used in conditional checks, e.g.,
# if (require(ggplot2)) {
#   print("ggplot2 is available and loaded.")
# } else {
#   print("ggplot2 is not available. Please install it.")
# }

# Example of a package that might not be installed, to demonstrate require()
# This will likely return FALSE and print the 'else' message
if (require(nonexistentpackage123, quietly = TRUE)) {
  print("nonexistentpackage123 is available and loaded.")
} else {
  print("nonexistentpackage123 is not available or could not be loaded.")
}

## [1] "nonexistentpackage123 is not available or could not be loaded."
```

Note: Once a package is loaded, its functions are available directly by name. If two loaded packages have functions with the same name, R will use the one from the package loaded *last*. You can always explicitly specify a function's package using `package_name::function_name()`. For example, `dplyr::filter()` if you want to ensure you're using `dplyr`'s filter function.

4. Finding and Exploring Packages The R community has created an enormous number of packages. Knowing where to find them and how to explore their functionalities is crucial.

- **CRAN (Comprehensive R Archive Network):** The official repository for R packages.
 - **Website:** <https://cran.r-project.org/web/packages/>
 - You can browse packages by name, task view (e.g., “Machine Learning”, “Time Series”), or author. Each package has a dedicated page with a description, downloads, and links to documentation.
- **Bioconductor:** A specialized repository for bioinformatics and computational biology packages.
 - **Website:** <https://www.bioconductor.org/>

- **GitHub:** Many developers host packages on GitHub before or instead of submitting to CRAN. You can install directly from GitHub using the `install_github()` function from the `devtools` package.
- **R Documentation:** Once a package is loaded, you can access its documentation:
 - `?function_name`: Opens the help page for a specific function.
 - `help(package = "package_name")`: Opens the main help page for the entire package.
 - `example(function_name)`: Runs the examples provided in the function’s help page.

Code Snippets:

```
print("--- Finding and Exploring Packages ---")

## [1] "--- Finding and Exploring Packages ---"
# Access help for a function (e.g., from base R)
# ?mean # Uncomment to open help page for 'mean'

# Access help for a function from a loaded package (e.g., stringr::str_length)
# ?str_length # Uncomment to open help page for 'str_length'

# View all functions/objects in a loaded package
# ls("package:stringr") # Uncomment to list all objects in stringr namespace

# Run examples from a function's help page
# example(str_length) # Uncomment to run examples for str_length
```

Expected (if uncommented): (Output would appear in R’s help viewer or console)

5. Commonly Used Packages Here are some essential R packages that are widely used in data science and analysis:

- **dplyr (part of tidyverse):** For data manipulation (filtering, selecting, arranging, summarizing, mutating data frames). It provides a consistent and intuitive set of “verbs” for data wrangling.
- **ggplot2 (part of tidyverse):** For data visualization. It implements a “grammar of graphics,” allowing you to build complex and beautiful plots layer by layer.
- **readr (part of tidyverse):** For fast and friendly reading of rectangular data (CSV, TSV, etc.). Often faster and more robust than base R `read.csv()`.
- **tidyr (part of tidyverse):** For “tidying” data, making it easier to work with. Functions like `pivot_wider()` and `pivot_longer()` are key for reshaping data.
- **purrr (part of tidyverse):** For functional programming, making it easier to work with lists and apply functions iteratively.
- **lubridate (part of tidyverse):** For working with dates and times.
- **stringr (part of tidyverse):** For consistent and convenient string manipulation.
- **forcats (part of tidyverse):** For working with factors (categorical variables).
- **data.table:** An extremely fast and memory-efficient package for data manipulation, particularly for very large datasets. It has a different syntax compared to `dplyr`.
- **caret:** For machine learning workflows (Classification And REgression Training). Provides a unified interface for many machine learning models.
- **Shiny:** For building interactive web applications directly from R.

The tidyverse: Many of these commonly used packages (`dplyr`, `ggplot2`, `readr`, `tidyr`, `purrr`, `stringr`, `forcats`, `lubridate`) are part of a larger ecosystem called the **tidyverse**. Installing `install.packages("tidyverse")` will install all of them at once. Loading `library(tidyverse)` will load the core tidyverse packages.

Code Snippets (Illustrating tidyverse usage):

```

# Install tidyverse (if you haven't already - this can take a few minutes)
# install.packages("tidyverse") # Uncomment to run

# Load the tidyverse (this loads dplyr, ggplot2, readr, etc.)
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v purrr      1.0.4
## v forcats    1.0.0      v readr      2.1.5
## v ggplot2    3.5.2      v tibble     3.2.1
## v lubridate  1.9.4      v tidyr      1.3.1
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

print("tidyverse loaded (dplyr, ggplot2, readr, etc. are now available).")

## [1] "tidyverse loaded (dplyr, ggplot2, readr, etc. are now available)."
```

Example using dplyr: Filtering a built-in dataset
mtcars is a built-in dataset in R

```

filtered_cars <- mtcars %>%
  filter(cyl == 4 & gear == 4) %>%
  select(mpg, hp, wt)

print("Filtered cars (using dplyr):")

## [1] "Filtered cars (using dplyr):"
```

```

print(head(filtered_cars))


##           mpg  hp  wt
## Datsun 710  22.8  93 2.320
## Merc 240D  24.4  62 3.190
## Merc 230   22.8  95 3.150
## Fiat 128   32.4  66 2.200
## Honda Civic 30.4  52 1.615
## Toyota Corolla 33.9 65 1.835
```

Example using ggplot2: Simple plot (won't display in text output)

```

ggplot(mtcars, aes(x = mpg, y = hp)) +
  geom_point() +
  labs(title = "MPG vs HP in mtcars") # Uncomment to create a plot in your R environment

```



This lesson provided a comprehensive introduction to R packages, explaining their importance, how to install and load them, and how to find their documentation. You also got an overview of some of the most widely used packages in the R ecosystem, particularly those within the **tidyverse**. Packages are the key to unlocking R's full potential for data analysis and beyond.

Module 1.4: Functions and Packages is now complete!

Next, we will begin **Phase 2: R Programming Fundamentals**, starting with **Module 2.1: Data Manipulation with dplyr and tidyr**.